

Durham Research Online

Deposited in DRO:

13 July 2009

Version of attached file:

Accepted Version

Peer-review status of attached file:

Peer-reviewed

Citation for published item:

Andreou, P. C. and Charalambous, C. and Martzoukos, S. H. (2008) 'Pricing and trading European options by combining artificial neural networks and parametric models with implied parameters.', *European journal of operational research.*, 185 (3). pp. 1415-1433.

Further information on publisher's website:

<http://dx.doi.org/10.1016/j.ejor.2005.03.081>

Publisher's copyright statement:

Additional information:

Use policy

The full-text may be used and/or reproduced, and given to third parties in any format or medium, without prior permission or charge, for personal research or study, educational, or not-for-profit purposes provided that:

- a full bibliographic reference is made to the original source
- a [link](#) is made to the metadata record in DRO
- the full-text is not changed in any way

The full-text must not be sold in any format or medium without the formal permission of the copyright holders.

Please consult the [full DRO policy](#) for further details.

Pricing and Trading European Options by Combining Artificial Neural Networks and Parametric Models with Implied Parameters

**Panayiotis C. Andreou¹, Chris Charalambous², and Spiros H.
Martzoukos^{*3}**

University of Cyprus – Department of Public and Business Administration

This version: March 2005

Keywords: Finance, neural networks, empirical option pricing

Acknowledgements: This work has been partly funded by the Hermes European Center of Excellence in Computational Finance and Economics, and a University of Cyprus grant for research in ANNs and Derivatives. We are thankful to Paul Lajbcygier for comments and discussions. We appreciate comments on earlier drafts by participants at the 9th Annual Conference of the Multinational Finance Society (Paphos, July 2002), the International Conference of Artificial Neural Networks (Madrid, August 2002), the International Workshop on Computational Management Science (Limassol, March 2003), the joint EURO/INFORMS Meeting (Istanbul, July 2003), the European Union Workshop series on Mathematical Optimization Models for Financial Institutions (Ayia Nappa, November 2003), and the FMA International/Europe Conference (Zurich, June 2004).

1,2,3 PhD Candidate, Professor of Management Science, and Assistant Professor of Finance respectively

* Corresponding author:

Spiros H. Martzoukos, Assistant Professor of Finance
University of Cyprus, Department of Public and Business Administration
P.O. Box 20537, CY 1678 Nicosia – Cyprus
Fax: +357-22 89 24 60, Tel.: +357-22 89 24 74
email: baspiros@ucy.ac.cy

Pricing and Trading European Options by Combining Artificial Neural Networks and Parametric Models with Implied Parameters

Abstract

We compare the ability of the parametric Black and Scholes, Corrado and Su models, and Artificial Neural Networks to price European call options on the S&P 500 using daily data for the period January 1998 to August 2001. We use several historical and implied parameter measures. Beyond the standard neural networks, in our analysis we include hybrid networks that incorporate information from the parametric models. Our results are significant and differ from previous literature. We show that the Black and Scholes based hybrid artificial neural network models outperform the standard neural networks and the parametric ones. We also investigate the economic significance of the best models using trading strategies (extended with the Chen and Johnson modified hedging approach). We find that there exist profitable opportunities even in the presence of transaction costs.

1. Introduction

In this paper we compare parametric option pricing models (OPMs) -- Black and Scholes (1973) (BS) and the semi-parametric Corrado and Su (1996) (CS) -- with several artificial neural network (ANN) configurations. We compare them with respect to pricing the S&P 500 European call options, and trading strategies are implemented in the presence of transaction costs.

Black and Scholes introduced in 1973 their milestone OPM. Despite the fact that BS and its variants are considered as the most prominent achievements in financial theory in the last three decades, empirical research has shown that the formula suffers from *systematic biases* (see Black and Scholes, 1975, MacBeth and Merville, 1980, Gultekin et al., 1982, Rubinstein, 1985, Bates, 1991 and 2003, Bakshi et al., 1997, Andersen et al., 2002, and Cont and Fonseca, 2002). The BS bias stems from the fact that the model has been developed under a set of simplified assumptions such as geometric Brownian motion of stock price movements, constant variance of the underlying returns, continuous trading on the underlying asset, constant interest rates, etc.

Post-BS research (e.g. stochastic volatility, jump-diffusion, stochastic interest rates, etc.) has not managed to either generalize all the assumptions of BS or provide results truly consistent with the observed market data. These models are often too complex to implement, have poor out-of-sample pricing performance and have implausible and sometimes inconsistent implied parameters (see Bakshi et al., 1997). This justifies the severe time endurance of BS¹. Together with the BS model, we also consider the semi-parametric CS model that allows for excess skewness and kurtosis, as a model that can proxy for other more complex parametric ones.

Nonparametric techniques such as *Artificial Neural Networks* are promising alternatives to the parametric OPMs. ANNs do not necessarily involve directly any financial theory because the option's price is estimated inductively using historical or implied input variables and option transactions data. Option-pricing functions are multivariate and highly nonlinear, so ANNs are desirable approximators of the *empirical option pricing function*. Parametric models describe a stationary nonlinear relationship between a theoretical option price and various variables. Since it is known that market participants change their option pricing attitudes from time to time (i.e. Rubinstein, 1985) a stationary model may fail to adjust to such rapidly changing market behavior (see also Cont and Fonseca, 2002, for evidence of noticeable variation in daily implied parameters). ANNs if frequently trained can

¹ According to Andersen et al., (2002), "*the option pricing formula associated with the Black and Scholes diffusion is routinely used to price European options, although it is known to produce systematic biases*".

adapt to changing market conditions, and can potentially correct the aforementioned BS bias (Hutchison et al., 1994, Lajbcygier et al., 1996, Garcia and Gencay, 2000, Yao and Tan, 2000).

Beyond the standard ANN target function we further examine the hybrid ANN target function suggested by Watson and Gupta (1996) and used for pricing options with ANNs in Lajbcygier et al. (1997). In the hybrid models the target function is the residual between the actual call market price and the parametric option price estimate. In previous studies the *standard steepest descent backpropagation* algorithm is (mostly) used for training the feedforward ANNs. It is shown in Charalambous (1992) that this learning algorithm is often unable to converge rapidly to the optimal solution. Here we utilize the modified Levenberg-Marquardt (LM) algorithm which is much more sophisticated and efficient in terms of time capacity and accuracy (Hagan and Menhaj, 1994). In contrast to most previous studies, thorough cross-validation allows us to use a different network configuration in different testing periods.

The data for this research come from two dominant world markets, the New York Stock Exchange (NYSE) for the S&P 500 equity index and the Chicago Board of Options Exchange (CBOE) for call option contracts, spanning a period from January 1998 to August 2001. To our knowledge, the resulting dataset is larger than the ones used in other published studies. We also (similarly to Rubinstein, 1985, Bates, 1996, Bakshi et al., 1997; see discussion in Bates, 2003) *reserve* option datapoints that in several ANN studies were dropped out of the analysis. Note that in order to check the robustness of the results we repeated the analysis using a reduced dataset following Hutchison et al. (1994). We examine more explanatory variables including historical, weighted average implied and pure implied parameters. Also, instead of constant maturity riskless interest rate, we use nonlinear interpolation for extracting a continuous rate according to each option's time to maturity.

Lastly, although previous researchers have exploited BS or ANNs, little has been reported for the case of CS² and nothing for the hybrid ANNs that use information derived by CS. To investigate the economic significance of the alternative option pricing approaches, trading strategies without and with the inclusion of transaction costs are utilized. These trading strategies are implemented with the standard delta-hedging values implied by each model, but also with the corrected values according to the (widely neglected) Chen and Johnson (1985) methodology.

In the following we first review the BS and CS models, and the standard and hybrid ANN model configuration. Then we discuss the dataset, the historical and

² An exception is the paper by Sami Vahamaa (2003) that examined the hedging performance of the CS model without including transaction costs.

implied parameter estimates we derive, and we define the parametric and ANN models according to the parameters used. Subsequently we review the numerical results with respect to the in- and out-of-sample pricing errors; and we discuss the economic significance of dynamic trading strategies both in the absence and in the presence of transaction costs. The final section concludes. In general, our results are novel and significant. We identify the best hybrid ANN models, and we provide evidence that (even in the presence of transaction costs), profitable trading opportunities still exist.

2. Option pricing: BS, CS and ANNs

2.1. The parametric models

The Black Scholes formula for European call options modified for dividend-paying underlying asset is:

$$c^{BS} = Se^{-\delta T} N(d_1) - Xe^{-rT} N(d_2), \quad (1)$$

where,

$$d_1 = \frac{\ln(S/X) + (r - \delta)T + (\sigma\sqrt{T})^2 / 2}{\sigma\sqrt{T}}, \quad (1.a)$$

and

$$d_2 = d_1 - \sigma\sqrt{T}. \quad (1.b)$$

c^{BS} \equiv premium paid for the European call option; S \equiv spot price of the underlying asset; X \equiv exercise price of the option; r \equiv continuously compounded riskless interest rate; δ \equiv continuous dividend yield paid by the underlying asset; T \equiv time left until the option expiration; σ^2 \equiv yearly variance rate of return for the underlying asset; $N(.)$ \equiv the standard normal cumulative distribution.

The standard deviation of continuous returns (σ) is the only variable in Equations 1.a and 1.b that cannot be directly observed in the *market*. For this study, we use both *historical* and *implied volatility* forecasts. For the *Historical Volatility* we use the past 60 days. The *Implied Volatility* (IVL) calculation involves solving Equation 1 iteratively for σ given the values of the observable c^{mrk} (the most recently observed market price of a call option), and the relevant values of S , X , T , r and δ . Contrary to historical volatility, IVL has desirable properties that make it

attractive to practitioners: it is forward looking, and avoids the assumption that past volatility will be repeated.

If BS is a well-specified model, then all IVLs on the same underlying asset should be the same, or at least deterministic functions of time. Unfortunately, many researchers have reported systematic biases. For example, Rubinstein (1985) has shown that IVL derived via BS as a function of the moneyness ratio (S/X) and time to expiration (T) often exhibits a **U** shape, the well known *volatility smile*. Bakshi et al. (1997) report that implicit stock returns' distributions are negatively skewed with more excess kurtosis than allowable in the BS lognormal distribution. This is why we usually refer to BS as being a misspecified model with an inherent source of bias (see also Latane and Rendleman, 1976, Bates, 1991, Canica and Figlewski, 1993, Bakshi et al., 2000, and Andersen et al., 2002). For the aforementioned reason we include in our analysis the Corrado and Su (1996) (see also the correction in Brown and Robinson, 2002) model that explicitly allows for excess skewness and kurtosis. The CS model is a semi-parametric model since it does not rely on specific assumptions about the underlying stochastic process. Corrado and Su define their model as:

$$c^{CS} = c^{BS} + \mu_3 Q_3 + (\mu_4 - 3)Q_4, \quad (2)$$

$$Q_3 = \frac{1}{3!} Se^{-\delta T} \sigma \sqrt{T} ((2\sigma \sqrt{T} - d_1)n(d_1) - \sigma^2 T N(d_1)), \quad (2.a)$$

$$Q_4 = \frac{1}{4!} Se^{-\delta T} \sigma \sqrt{T} ((d_1^2 - 1 - 3\sigma \sqrt{T}(d_1 - \sigma \sqrt{T}))n(d_1) + \sigma^3 T^{3/2} N(d_1)), \quad (2.b)$$

$$n(z) = \frac{1}{\sqrt{2\pi}} \exp(-z^2 / 2), \quad (2.c)$$

where c^{BS} is the BS value for the European call option adjusted for dividends, and μ_3 and μ_4 are the coefficients of skewness and kurtosis of the returns.

2.2. Neural networks

A Neural Network is a collection of interconnected simple processing elements structured in successive layers and can be depicted as a network of arcs/connections and nodes/neurons. Fig. 1 depicts a fully-connected ANN architecture similar to the one applied in this study. This network has three layers: an *input* layer with N input variables, a *hidden* layer with H neurons, and a single

neuron *output* layer. Each connection is associated with a *weight*, w_{ki} , and a *bias*, b_k , in the hidden layer and a *weight*, v_k , and a *bias*, v_0 , for the output layer ($k = 1, 2, \dots, H$, $i = 1, 2, \dots, N$). A particular neuron node is composed of: *i*) the vector of *input signals*, *ii*) the *vector weights* and the associated *bias*, *iii*) the *neuron* itself that *sums* the product of the input signal with the corresponding weights and bias, and finally, *iv*) the *neuron transfer function*. In addition, the outputs of the hidden layer ($y_1^{(1)}, y_2^{(1)} \dots y_H^{(1)}$) are the inputs for the output layer. Since we want to approximate the market options pricing function, ANNs operate as a non-linear regression tool:

$$Y = G(\tilde{x}) + \varepsilon_{ANN}, \quad (3)$$

that maps the unknown relation, $G(\cdot)$, between the input variable vector, $\tilde{x} = [x_1, x_2, \dots, x_N]$, the target function, Y , and the error term, ε_{ANN} . Inputs are set up in feature vectors, $\tilde{x}_q = [x_{1q}, x_{2q}, \dots, x_{Nq}]$ for which there is an associated and known target, $Y \equiv t_q$ (in our case, $t_q \equiv c_q^{mrk} / X_q$), with $q = 1, 2, \dots, P$, where P is the number of the available sample features. According to Fig. 1, the operation carried out for estimating output y (in our case, $y_q \equiv c_q^{ANN} / X_q$), is the following:

$$y = f_0[v_0 + \sum_{k=1}^H v_k f_H(b_k + \sum_{i=1}^N w_{ki} x_i)]. \quad (4)$$

[Figure 1 here]

For the purpose of this study, the hidden layer always uses the hyperbolic tangent sigmoid transfer function, while the output layer uses a linear transfer function. In addition, ANN architectures with only one hidden layer are considered since they operate as a nonlinear regression tool and can be trained to approximate most functions arbitrarily well (Cybenko, 1989). High accuracy can be obtained by including enough processing nodes in the hidden layer.

To train the ANNs, we utilized the modified LM algorithm. According to LM, the weights and the biases of the network are updated in such a way so as to minimize the following sum of squares performance function:

$$F(W) = \sum_{q=1}^P e_q^2 \equiv \sum_{q=1}^P (y_q - t_q)^2 \equiv \sum_{q=1}^P (f_0[v_0 + \sum_{k=1}^H v_k f_H(b_k + \sum_{i=1}^N w_{ki} x_{iq})] - t_q)^2, \quad (5)$$

where, W is an n -dimensional column vector containing the weights and biases: $W = [b_1, \dots, b_H, w_{11}, \dots, w_{HN}, v_0, \dots, v_H]^T$. Then, at each iteration τ of LM, the weights vector W is updated as follows:

$$W_{\tau+1} = W_{\tau} - [J^T(W_{\tau})J(W_{\tau}) + \mu_{\tau}I]^{-1}J^T(W_{\tau})e(W_{\tau}), \quad (6)$$

where I is an $n \times n$ identity matrix, $J(W)$ is the $P \times n$ Jacobian matrix of the P -dimensional output error column vector $e(W)$, and μ_{τ} is like a learning parameter that is adjusted in each iteration in order to secure convergence. Further technical details about the implementation of LM can be found in Hagan and Menhaj (1994) and Hagan et al. (1996). In addition to the standard use of ANNs where $t_q \equiv c_q^{mrk} / X_q$, we also try hybrid ANNs in which the target function is the residual between the actual call market price and the BS or CS call option estimation:

$$t_q \equiv c_q^{mrk} / X_q - c_q^k / X_q, \quad (7)$$

with k defining inputs from a parametric model. To avoid neuron saturation, we scale input variables using the *mean-variance* transformation (z-score) defined as follows:

$$\tilde{z}_i = (\tilde{x}_i - \mu_i) / s_i, \quad (8)$$

where \tilde{x}_i is the vector containing all of the available observations related to a certain input/output variable for a specific training period, μ_i is the mean and s_i the standard deviation of this vector. Moreover, we also utilize the network initialization technique proposed by Nguyen and Windrow (see Hagan et al., 1996) that generates initial weights and bias values for a nonlinear transfer function so that the active regions of the layer's neurons are distributed roughly evenly over the input space.

In this study for each input variable set of each training sample, all the available networks having two to ten hidden neurons are cross-validated (in total nine). Moreover, since the initial network weights affect the final network performance, for a specific number of hidden neurons the network is initialized, trained and validated many times. Each network is estimated and optimized using the Mean Square Error (MSE) criterion shown in Equation 5 for no more than two-hundred iterations. The dataset is divided into three sub-sets. The first is the *training (estimation) set*. The second is the *validation set* where the ANN model's error

is monitored and the optimal number of hidden neurons and their weights are defined, via an early stopping procedure (MSE fails to decrease in 10 consecutive iterations). Given the optimal ANN structure, its pricing capability is tested in a third separate *testing dataset*.

3. Data, parameter estimates (historical and implied), and model implementation

Our dataset covers the period January 1998 to August 2001. To our knowledge, the resulting dataset is larger than the one used in other published studies and reserves option data points that in most of the previous studies were dropped out of the analysis. After implementing the filtering rules, our dataset consists of 76,401 data points, with an average of 35,000 data points per (overlapping rolling training-validation-testing) sub-period (see Fig. 2). Hutchison et al. (1994) have an average of 6,246 data points per sub-period. Lajbcygier et al. (1996) include 3,308 data points, Yao et al. (2000) include 17,790 data points, and Schittenkopf and Dorffner (2001) include 33,633 data points. The S&P 500 Index call options are considered because this option market is extremely liquid and one of the most popular index options traded on the CBOE. This market is the closest to the theoretical setting of the parametric models. Along with the index, we have collected a daily dividend yield, δ , provided online by Datastream.

3.1. Observed and historically estimated parameters

Moneyiness Ratio (S/X): The moneyiness ratio may explicitly allow the ANNs to learn the moneyiness bias associated with the BS (see also Garcia and Gencay, 2000). The dividend adjusted moneyiness ratio $(Se^{-\delta T})/X$ is used in this study with ANNs because it is more informative. The simple moneyiness ratio S/X is used in order to tabulate results as in Hutchison et al. (1994). We adopt the following terminology: *very deep out of the money* (VDOTM) when $S/X < 0.85$, *deep out the money* (DOTM) when $0.85 \leq S/X < 0.90$, *out the money* (OTM) when $0.90 \leq S/X < 0.95$, *just out the money* (JOTM) when $0.95 \leq S/X < 0.99$, *at the money* (ATM) when $0.99 \leq S/X < 1.01$, *just in the money* (JITM) when $1.01 \leq S/X < 1.05$, *in the money* (ITM) when $1.05 \leq S/X < 1.10$, *deep in the money* (DITM) when $1.10 \leq S/X < 1.35$, and *very deep in the money* (VDITM) when $S/X \geq 1.35$.

Time to maturity (T): For each option contract, trading days are computed assuming 252 days in a year. In terms of time length, an option contract is classified

as *short term maturity* when its maturity is less than 60 days, as *medium term maturity* when its maturity is between 60 and 180 days and as *long term maturity* when it has maturity longer than (or equal to) 180 days.

Riskless interest rate (r): Most of the studies use 90-day T-bill rates (or similar when this is unavailable) as approximation of the interest rate. We use nonlinear cubic spline interpolation for matching each option contract with a continuous interest rate, r , that corresponds to the option's maturity, by utilizing the 3-month, 6-month and one-year T-bill rates collected from the U.S. Federal Reserve Bank Statistical Releases.

Historical Volatilities (σ): The 60-day historical volatility is calculated using all the past 60 log-relative index returns and is symbolized as σ_{60} .

CBOE VIX Volatility Index: It was developed by CBOE in 1993 and is a measure of the volatility of the S&P 100 Index³. VIX is calculated by taking the weighted average of the implied volatilities of eight S&P 100 Index call and put options with an average time to maturity of 30 days. This volatility measure can only be used with BS and is symbolized as σ_{vix}^{BS} .

Skewness and Kurtosis: The 60-day skewness ($\mu_{3,60}^{CS}$) and kurtosis ($\mu_{4,60}^{CS}$) needed for the CS model are approximated from the sixty most recent log-returns of the S&P 500.

3.2. Implied parameters

We adopt the Whaley's (1982) simultaneous equation procedure to minimize a price deviation function with respect to the unobserved parameters. As with Bates (1991), market option prices (c^{mrk}) are assumed to be the corresponding model prices (c^k , k defining input from a parametric model) plus a random additive disturbance term. For any option set of size N_t (N_t refers to the number of different call option transaction datapoints available on a specific day), the difference:

$$\varepsilon_{N_t}^k = c_{N_t}^{mrk} - c_{N_t}^k \quad (9)$$

between the market and the model value of a certain option is a function of the values taken by the unknown parameters. To find optimal implied parameter values we solve an unconstrained optimization problem that has the following form:

³ The S&P 100 Index and S&P 500 Index exhibit 30 day log-return average correlations for the period January 1998 to August 2002 of about 0.98.

$$SSE(t) = \min_{\theta^k} \sum_{l=1}^{N_t} (\varepsilon_l^k)^2, \quad (10)$$

where t represents the time instance, and θ^k the unknown parameters associated with a specific parametric OPM ($\theta^{BS} = \{\sigma\}$, $\theta^{CS} = \{\sigma, \mu_3, \mu_4\}$). The SSE is minimized via a non-linear least squares optimization based on the LM algorithm. To minimize the possibility to obtain implied parameters that correspond to a local minimum of the error surface (see also Bates, 1991, and Bakshi et al., 1997), with each model we use three different starting values for the unknown parameters based on reported average values in Corrado and Su (1996).

A difference of our approach compared to previous studies is that the above minimization procedure is used daily to derive four different sets of implied parameters for each parametric model. The first optimization is performed by including all available options transaction data in a day to obtain *daily average* implied structural parameters. Alternatively, for a certain day we minimize the SSE of Equation 10 by fitting the BS and CS for options that share the same maturity date as long as four different available call options exist. We thus get *daily average per maturity* parameters. In a third step, for every maturity each available option contract is grouped with its three nearest options in terms of the moneyness ratio in order to minimize the above SSE function, deriving thus parameters *average per the 4 closest* contracts; such estimates are ignored in previous research. We finally calibrate the implied structural parameters, by focusing on the Brownian volatility for each contract so as to drive the residual error to zero or to a negligible value. In the case of BS this is quite simple and we can easily obtain a *contract specific* volatility estimate. For CS we need three structural parameters, so for each call option we minimize Equation 10 with respect to the Brownian volatility after fixing the skewness and kurtosis coefficients to the values obtained from the previous procedure that gave the *average per the 4 closest* implied parameters. Two kinds of constraints are included in the optimization process for practical reasons: nonnegative implied volatility parameters are obtained by using an exponential transformation; and the skewness of CS⁴ is permitted to vary in the range $[-10, 5]$ whereas kurtosis is constrained to be less than 30. Unlike previous studies, we

⁴ If not somehow constrained, skewness and kurtosis can take implausible values (i.e. Bates, 1991) due to model overfitting that will lead to enormous pricing errors on the next day (especially for deep in the money options). In our case these constraints were binding in less than 2% of the whole dataset.

include contract specific implied parameters since these are widely used by market practitioners (i.e. Bakshi et al., 1997, pg. 2019).

For notational reasons, implied parameters obtained from the first step are denoted by the subscript av , from the second step by the subscript avT , from the third step by the subscript $avT4$, and from the fourth step by the subscript con . The four different implied BS volatility estimates are symbolized as: σ_j^{BS} , $j = \{av, avT, avT4, con\}$, whilst the four different sets of CS parameters as: $\{\sigma_j^{CS}, \mu_{3,j}^{CS}, \mu_{4,j}^{CS}\}$. For pricing and trading reasons at time instant t , the implied structural parameters derived at day $t-1$ are used together with all other needed information (S , T , X , r , and δ).

It is known that ANN input variables should be presented in a way that maximizes their information content. When we price options, the parametric OPM formulas adjust those values to represent the appropriate value that corresponds to an option's expiration period. According to this rationale, volatility measures for use with the ANNs are transformed by multiplying each of the yearly volatility forecast with the square root of each option's time to maturity ($\tilde{\sigma}_j = \sigma_j \sqrt{T}$, where $j = \{60, vix, av, avT, avT4, con\}$). We denote these volatility measures as $\tilde{\sigma}_j^{BS}$ and $\tilde{\sigma}_j^{CS}$; and we name them as *maturity (or expiration) adjusted volatilities*. Additionally, for the case of CS, skewness $\mu_{3,j}^{CS}$, $j = \{60, av, avT, avT4, con\}$, is transformed by multiplication with Q_3 that represents the marginal effect of nonnormal skewness. Similarly, $\mu_{4,j}^{CS}$ is multiplied with Q_4 . We denote these adjusted parameters as $\tilde{\mu}_{3,j}^{CS}$ (adjusted skewness), and $\tilde{\mu}_{4,j}^{CS}$ (adjusted kurtosis).

3.3. Output variables, filtering and processing

The BS (c_q^{BS}) and CS (c_q^{CS}) outputs, are used as an estimate for the market call option, c_q^{mrk} . For training ANNs, the call standardized by the striking price, c_q^{mrk} / X_q , is used as the target function to be approximated. In addition, we implement the hybrid structure where the target function represents the pricing error between the option's market price and the parametric models estimate, $c_q^{mrk} / X_q - c_q^k / X_q$.

Before filtering, more than 100,000 observations were included for the period January 1998 – August 2001. The filtering rules we adopt are: *i) eliminate an*

observation if the call contract price, $c_{m,t}^{mrk}$, m defining each traded contract, is greater than the underlying asset value, S_t ; ii) exclude an observation if the call moneyness ratio is larger than unity, $S_t/X_m > 1$, and the call price, $c_{m,t}^{mrk}$, is less than its lower bound, $S_t e^{-\delta_{m,t} T_{m,t}} - X_m e^{-r_{m,t} T_{m,t}}$; iii) eliminate all the options observations with time to maturity less than 6 trading days. The latter filtering rule is adopted to avoid extreme option prices that are observed due to potential illiquidity problems; iv) price quotes lower than 0.5 index points are not included; v) maturities with less than four call option observations are also eliminated, vi) in addition, to remove impact from thin trading we eliminate observations according to the following rule: eliminate an observation if the $c_{m,t}^{mrk}$ is equal to $c_{m,t-1}^{mrk}$ and if the open interest for these days stays unchanged and if the underlying asset S has changed.

[Table 1, here]

Our final dataset consists of 76,401 datapoints. Table 1 exhibits some of the properties of our sample tabulated according to moneyness ratio and time to maturity forming 27 different moneyness/maturity classes. We provide the average values for c^{mrk} and σ_{con}^{BS} , and the number of observations within each moneyness and maturity class. The implied volatility, σ_{con}^{BS} , presents a non-flat moneyness structure when fixing the time to maturity and vice versa revealing the bias associated with BS. Moreover, we should notice that DITM and VDITM options dominate in number of datapoints all other classes, so unlike studies that ignore these options we choose to include them in the dataset. For the training sub-periods, the observations vary between: 19,852-22,545; for the validation sub-periods between: 10,372-10,916; and for the testing sub-periods between: 3,797-4,264.

In order to check the robustness of the results, in addition to the *full* dataset just described, we repeat the analysis using a *reduced* dataset. In this reduced dataset we follow Hutchison et al. (1994), and we neither use long maturity (longer than 180 trading days) options, nor the VDOTM ($S/X < 0.85$) or the VDITM ($S/X \geq 1.35$) options. The excluded observations (because of considerations of thin trading) comprise about 21% of the full dataset resulting in a total of 60,402 observations. The training-validation-testing splitting dates are the same as in the original dataset. For the training sub-periods, the observations vary between: 15,851-18,053; for the validation sub-periods: 7,728-9,638; and for the testing sub-periods: 2,689-3,983. To be consistent with Hutchison et al. (1994), in using the reduced dataset we

retrain the ANNs. Our discussion will focus on the full dataset. In order to save space, we will only show selected results using the reduced dataset.

3.4. Validation and testing, and pricing performance measures

Since a practitioner is faced with time-series data, it was decided to partition the available data into training, validation and testing datasets using a chronological manner, and via a rolling-forward procedure. Our dataset is divided into ten different overlapping training (Tr) and validation (Vd) sets, each followed by separate and non-overlapping testing (Ts) sets as exhibited by Fig. 2. The ten sequential testing sub-periods cover the last 25 months of the complete dataset.

[Figure 2, here]

There are M available call option contracts, for each of which there exist Ξ_m observations taken in consecutive time instances t , resulting in a total of P ($P = \sum_{m=1}^M \Xi_m$) available call option datapoints. To determine the pricing accuracy of each model's estimates c^k (k defining the model), we examine the Root Mean Square Error (RMSE) and the Mean Absolute Error (MAE):

$$RMSE = \sqrt{(1/p) \sum_{v=1}^p (c_v^{mrk} - c_v^k)^2}, \quad (11)$$

$$MAE = (1/p) \sum_{v=1}^p |c_v^{mrk} - c_v^k|, \quad (12)$$

where p indicates the number of observations. The error measures are computed for an aggregate testing period ($AggTs$) with 39,831 datapoints by pooling together the pricing estimates of all ten testing periods. For $AggTs$ we also compute the Median of the Absolute Error (MeAE). Of course, since ANNs are effectively optimized with respect to the mean square error, the out-of-sample pricing performance should be similarly based on RMSE and in a lesser degree on MAE and MeAE.

3.5. The alternative BS, CS and ANN models

With the BS models we use as input S , X , T , r , δ , and any of the six different volatility measures: σ_{60} , σ_{vix}^{BS} , σ_{av}^{BS} , σ_{avT}^{BS} , σ_{avT4}^{BS} and σ_{con}^{BS} . Using P in the superscript

to denote the parametric version of BS, the six different models are symbolized as: BS_{60}^P , BS_{vix}^P , BS_{av}^P , BS_{avT}^P , BS_{avT4}^P , and BS_{con}^P . In a similar way there are five different CS models according to the kind of parameters used: CS_{60}^P , CS_{av}^P , CS_{avT}^P , CS_{avT4}^P , and CS_{con}^P .

With ANNs, we also use three standard input variables/parameters: $(Se^{-\delta T})/X$, T and r . Additional input parameters depend on the parametric model considered. There are six ANN models that use as an additional input the above BS volatility measures to map the standard target function c^{mrk}/X . There are six more versions that utilize the maturity adjusted parameters. Each of the previous input parameter sets is also used with the hybrid target function. The ANNs that use the untransformed BS volatility forecast are denoted by N in the superscript, the transformed versions by N^* , while the corresponding hybrid versions by Nh and Nh^* respectively. For instance, BS_{con}^N (BS_{con}^{Nh}) is the ANN model that uses as additional input σ_{con}^{BS} and maps the standard (hybrid) target function, whilst BS_{con}^{N*} (BS_{con}^{Nh*}) the ANN model that uses as additional input $\tilde{\sigma}_{con}^{BS}$ and maps the standard (hybrid) target function. In total there are 24 different versions of ANNs related to the BS and 20 related to the CS model.

5. Pricing results and discussion

We briefly review the observed in-sample fit of the parametric models as well as the in-sample characteristics of the various implied parameters. Then we discuss the out-of-sample performance of the alternative OPMs. When we do not explicitly refer to the dataset, we imply the full one. The insights derived were not affected by the choice of dataset. When noteworthy differences exist, we state them explicitly.

5.1. BS and CS in-sample fitting performance and implied parameters

Based on our (not reported in detail for brevity) statistics for the whole period (1998-2001) we have observed that CS is producing smaller fitting errors than the BS. The contract specific fitting procedure reduces the fitting errors so as to almost eliminate the residuals and obtain fully calibrated implied parameters. The in sample RMSE measures using the overall average set of implied parameters (av), the average per maturity (avT), and the closest four contracts ($avT4$), are: 11.63, 11.31, and 7.00

for the BS model; and 9.52, 8.52, and 5.35 for the CS model⁵. From unreported statistics we can also attest that the S&P 500 average σ_{con}^{BS} in 1998 was about 33%, in 1999 about 30%, in 2000 about 26% and in 2001 about 27%. It seems that the in-sample fitting error of the models (diminishing in time) is positively correlated with the market volatility.

We can also provide some statistics about the implied parameter values for the whole period. The Brownian volatility varies between 22% and 30% in BS and between 27% and 31% in CS. For the BS model, the average implied volatility (σ_{av}^{BS}) estimates are smaller in magnitude (both in mean and in median values) from the contract specific implied volatility, σ_{con}^{BS} , although similar volatility estimates do not necessarily lead to similar pricing and hedging values (Bakshi et al., 1997). Regarding implied skewness and kurtosis, the implicit distributions are negatively skewed with excess kurtosis in almost all days, something that is probably attributed to the crash fears of the market participants after the Black Monday of 1987. Implied average skewness does not change significantly (from -1.19 to -1.20) if we move from $\{av\}$ to $\{avT\}$ but there is a shift in implied average kurtosis (from 6.91 to 6.19).

5.2. Out-of-sample pricing results

Table 2, exhibits the performance of all parametric and ANN models considered in this study in terms of RMSE, MAE and MeAE for the *AggTs* (aggregate) period. In Table 3 we tabulate statistics for a pairwise comparison of the (statistical significance of) pricing performance of a selection of models. Since the ten testing periods are disjoint and because we have pricing estimates coming from different OPMs we can assume (similarly to Hutchison et al, 1994 and Schittenkopf and Dorffner, 2001) that the pricing errors are independent and standard t-test can be applied. Similarly to the previous authors we need to report that these tests should be interpreted with caution. The upper diagonal of Table 3 reports the *t*-values taken by a two-tail matched-pair test about the MAE of the alternative models whilst the lower diagonal exhibits the two-tail matched-pair *t*-test values about the MSE of the compared OPMs. Table 4 provides (as a robustness check) the performance of the models when using the reduced dataset.

[Table 2, 3 and 4, here]

⁵ The RMSE for CS in the fourth step (*con*) is 1.82 (caused by a tiny part of the dataset less than 0.1%) due to binding constraints on skewness and kurtosis. For this step, the MeAE is more appropriate, and is effectively zero. The RMSE and the MeAE for BS in the fourth step are effectively zero.

By looking at Tables 2 and 4 we can see that the use of implied instead of historical parameters improves performance, both for parametric and ANN models (in both datasets). Note that the 60-day historical volatility performed better than VIX with the parametric BS model, but the VIX volatility measure performed better with the ANN models. Using time adjusted parameters in the ANNs or using contract specific parameters $\{avT4, con\}$ usually improves performance. The combination of time adjusted parameters and contract specific parameters always provided the best model within each class of ANNs (standard or hybrid, BS or CS based) in both datasets.

In comparing the parametric models and again looking at Tables 2 and 4, it is noteworthy that CS outperforms BS when average implied parameters are used. BS still works better with contract specific parameters. The overall best among the parametric models is the contract specific BS model. In other more complex parametric models that include jumps and stochastic volatility components (i.e. Bakshi et al., 1997), deriving implied parameters may lead to model overfitting. The contract specific approach we adopt in this study seems not to lead to model overfitting, retaining thus good out-of-sample properties. For the ANN models, the CS based may outperform the BS based in some cases, but when the best combinations are used (time adjusted parameters and contract specific parameters), the best model always is BS based in both the standard and hybrid networks.

In comparing the parametric models with the standard ANNs, in the full dataset the ANNs never outperform the equivalent parametric ones. Apparently, the standard ANNs cannot perform well in the extreme data regions. In the reduced dataset (see Table 4), we observe the opposite since the standard ANNs always outperform the equivalent parametric ones.

In comparing the hybrid with the standard ANNs, in the full dataset the hybrid are always better. In the reduced dataset this may not always be the case, but the best combinations (time adjusted parameters and contract specific parameters) give as the best model always a hybrid one.

In both the full and the reduced dataset, the hybrid always outperform the equivalent parametric ones. Finally, in both the full and the reduced dataset, the overall best model is the BS based hybrid with time adjusted and contract specific volatility.

From Table 3, we can confirm the statistical significance of the best models. The comparative results we discuss with tests using the full dataset, and they also hold for the reduced dataset (statistics not reported for brevity). We can see that

BS_{con}^{Nh*} outperforms all other models. Specifically, BS_{con}^{Nh*} is producing a RMSE equal to 6.01 and a MAE equal to 2.61, pricing measures that are smaller than any other model at the 5% significance level.

The BS based hybrid ANNs even with historical or the VIX volatility measure are considerably better than the equivalent parametric alternatives at a statistically significant level. Specifically, BS_{60}^{Nh*} is producing 1.23 (1.25) times smaller MSE (MAE) compared to BS_{60}^P . Also BS_{vix}^{Nh*} produces 1.52 (1.90) times smaller MSE (MAE) compared to BS_{vix}^P .

Comparing the out-of-sample pricing performance of BS_{con}^{Nh*} to CS_{con}^{Nh*} we observe that the extra ANN flexibility of the latter due to the two additional input parameters does not lead to increased accuracy. The BS_{con}^{Nh*} is better than the CS_{con}^{Nh*} model at 1% significance level.

We can similarly see the statistical significance of the superiority of the BS based models with contract specific volatility versus the equivalent CS based models (both parametric and hybrid); and the superiority of the models using the implied volatility versus the equivalent ones using the historical volatility measures.

5.3. Other statistics

We tabulate in Table 5 the MSE of a selective (but representative) choice of models, according to the various moneyness and maturity classes for the aggregate (*AggTs*) period. We demonstrate results for the two best performing parametric models which serve as benchmark (BS_{con}^P , CS_{con}^P) and the two best performing (in their respective class) hybrid ANN models (BS_{con}^{Nh*} , CS_{con}^{Nh*}). We also demonstrate results for the reduced dataset (BS_{con}^{Nh*} , CS_{con}^{Nh*}). The relevant information for the parametric models in the reduced dataset can be taken from the information concerning the full if we ignore the long maturities, and the VDOTM and the VDITM classes. Very briefly, what can be seen is that BS_{con}^P has a smaller RMSE in all data classes compared to CS_{con}^P . The same holds for BS_{con}^{Nh*} over CS_{con}^{Nh*} . If we compare the BS and CS based hybrid models with the equivalent parametric ones, the hybrid ANN models rarely underperform the parametric ones, and they do so only in some classes far away from ATM. This we attribute to the scarcity of such call option datapoints in the training samples compared to other moneyness and maturity classes.

[Table 5, here]

We should finally comment on the complexity of each neural network configuration. Since we have a constant number of inputs within each model class, the larger the number of hidden neurons the more complex the ANN model architecture, and the more complex the target function to be approximated. Firstly, we observe that the number of hidden neurons changes significantly between sub-periods. This contradicts many previous studies that employ the assumption that the market's options pricing mechanism is the same for all periods examined and that a constant ANN structure is sufficient. Secondly, the standard target function is more complex compared to the hybrid one, hence this hybrid category of networks can perform better in out-of-sample pricing. Thus, it is not surprising that the best performing ANN model, BS_{con}^{Nh*} , demonstrates the simplest structure with an average of 3.2 hidden layer neurons, compared to the 8 hidden layer neurons in the case of the equivalent standard ANN (BS_{con}^{N*}). Similarly for the CS-based ANNs, we have 4.9 (for CS_{con}^{Nh*}) and 7.7 (for CS_{con}^{N*}) hidden layer neurons respectively. Similar network complexities (not reported) were observed in the reduced dataset.

6. Delta neutral trading strategies

We now investigate the economic significance of the best performing models in options trading. In order to save space we discuss the parametric versions of BS and CS which are usually the benchmark, and the hybrid ANN models which provided the overall best performance. Other studies usually restrict their analysis only to a hedging investigation of various alternative OPM models (i.e. Hutchison et al., 1994, Garcia and Gencay, 2000, Schittenkopf and Dorffner, 2001) and avoid exploiting trading strategies. It is known from previous studies that the best OPM in terms of out-of-sample pricing performance does not always prove to be the best solution when we consider delta hedging, since ANNs are optimized based on a pricing error criterion. Instead, and following the spirit of Black and Scholes (1972), Galai (1977), and Whaley (1982), we investigate the economic significance of the OPMs by implementing trading strategies. "A model that consistently achieves to identify mispriced options and within a time period produces an amount of trading profits will always be preferred by a practitioner" (Black and Scholes, 1972). The trading profitability that we will document, indirectly also hints to potential option market inefficiencies, although testing market efficiency is beyond the scope of our study. We implement trading strategies based on single instrument hedging, as for example in Bakshi et al. (1997). In addition, we consider various levels of transaction costs,

and we focus on dynamic strategies that are cost-effective. We later extend the analysis by implementing a modified approach for trading using hedging ratios obtained via the (widely neglected) Chen and Johnson (1985) method. To our knowledge, this is the first effort to validate this modified trading strategy using both parametric and ANN OPMs.

In the trading strategy we implement, we create portfolios by buying (selling) options undervalued (overvalued) relative to a model's prediction and taking a delta hedging position in the underlying asset. This (single-instrument) delta hedging follows the no-arbitrage strategy of Black and Scholes (1973), where a portfolio including a short (long) position in a call is hedged via a long (short) position in the underlying asset, and the hedged portfolio rebalancing takes place in discrete time intervals (in an optimal manner, not necessarily daily). At time t , if according to the model the m^{th} call option contract is overvalued (undervalued) relative to its market value, $c_{m,t}^{mrk}$, we go short (long) in this contract and we go long (short) in $\Delta_{m,t}^k$ "index shares"⁶, where k denotes the relevant model. Then we invest the residual, $B_{m,t}$, in a riskless bond. Note that $\Delta_{m,t}^k$ is the partial derivative of the option price with respect to the underlying asset, $\partial c_{m,t}^k / \partial S_t$, depending on the OPM under consideration. $\Delta_{m,t}^{ANN}$ can be calculated by differentiating Equation 4 via the chain rule. The expression for $\Delta_{m,t}^{BS}$ is $e^{-\delta T} N(d_1)$ and is derived from Equation 2.1. The expression for $\Delta_{m,t}^{CS}$ includes $\Delta_{m,t}^{BS}$ and is:

$$\Delta_{m,t}^{CS} = \Delta_{m,t}^{BS} + \mu_3 \Phi_3 + (\mu_4 - 3) \Phi_4, \quad (13)$$

where $\Phi_3 = \frac{\partial Q_3}{\partial S}$ and $\Phi_4 = \frac{\partial Q_4}{\partial S}$ are given below:

$$\Phi_3 = \frac{1}{3!} e^{-\delta T} ((\sigma\sqrt{T})^3 N(d_1) + n(d_1)[3(\sigma\sqrt{T})^2 - 3d_1\sigma\sqrt{T} + (d_1)^2 - 1]), \quad (13.a)$$

$$\begin{aligned} \Phi_4 = \frac{1}{4!} e^{-\delta T} ((\sigma\sqrt{T})^4 N(d_1) + 4n(d_1)((\sigma\sqrt{T})^3 - 6(d_1)n(d_1)(\sigma\sqrt{T})^2 - 4n(d_1)(\sigma\sqrt{T}) + \\ 4(d_1)^2 n(d_1)(\sigma\sqrt{T}) + 3(d_1)n(d_1) - (d_1)^3 n(d_1)). \end{aligned} \quad (13.b)$$

In general we avoid a *naive* (expensive) trading strategy with daily rebalancing, since in the presence of transaction costs this would become prohibitively expensive.

⁶ Similarly to Bakshi et al. (1997) we assume that the spot S&P 500 index is a traded security.

Instead, the position is held *as long as* the call is undervalued (overvalued) without necessarily daily rebalancing. Then the position is liquidated and the profit or loss is computed, tabulated separately and a new position is generated according to the prevailing conditions in the options market. This procedure is carried out for all contracts included in the dataset. We rebalance our position in the underlying asset to keep the appropriate hedge ratio. Rebalanced positions in the index, $V_{m,t+\Delta t}$, and the bond, $B_{m,t+\Delta t}$, are according to:

$$V_{m,t+\Delta t} = \pm S_{t+\Delta t} (\Delta_{m,t+\Delta t} - \Delta_{m,t}), \text{ and } B_{m,t+\Delta t} = B_{m,t} e^{r\Delta t} + V_{m,t+\Delta t}, \quad (14)$$

where the positive sign is considered when we treat undervalued and the negative sign when we treat overvalued options. Note that in all trading strategies, when we need to invest money we borrow and pay the riskless rate; similarly we do for as long as a strategy provides losses. Thus, when we present profits they are always above the dollar return on the riskless rate.

Computed statistics include the total profit or loss (P&L), the number of trades (# Trades), the total profit or loss at 0.2% transaction costs, P&L (0.2%), and 0.4% transaction costs, P&L (0.4%). The (proportional) transaction costs are paid for both positions (in the call option and in the “index shares”)⁷. We also implement strategies with enhanced cost-effectiveness by ignoring trades that involve call options whose absolute percentage mispricing error, $|c^k - c^{mrk}|/c^k$, is less than a mispricing margin $d = 15\%$, found as P&L ($d = 15\%$). In addition, for these strategies, we also calculate P&L under aggregate transaction costs for the “index shares”. With such aggregation, transactions in the underlying assets are paid on the net (aggregate) exposure of $V_{m,t+\Delta t}$ and not on each position individually. Under this strategy, we expect additional cost savings that may provide profits even at rather high transaction cost levels. We use the prefix *Agg.* in front of P&L to indicate this strategy. The following observations refer to the full dataset, but they also hold for the reduced one (unreported due to brevity considerations).

The results for the parametric BS and CS models are tabulated in Panel A of Tables 6 and 7 respectively. We observe that all models before transaction costs produce significant profits, implying that both BS and CS can successfully identify mispriced options. Within BS models the magnitude of P&L is larger for BS_{con}^P that

⁷ For example, assume that the index is at 1300 and a call option has a market price equal to 25 index points and a delta value of 0.60. Under 0.4% transaction costs the total commissions paid (for a single trade) will be 3.22 index points. In the *AggT*s period the S&P 500 was in a range from about 1100 to 1500. This level of transaction costs is low but attainable by professional traders and market makers.

employs a more *sophisticated* implied volatility forecast. Note though that the more sophisticated volatility forecast that is used with BS, the larger the number of trades. So, when 0.2% transaction costs are taken into consideration, all models produce significant losses and the previous profit dominance of BS_{con}^P over BS_{60}^P reverts because the latter model incurs less transaction costs (since it engages in a smaller number of trades). Similar results hold for the CS models although CS_{avT4}^P generates slightly higher profits compared to CS_{con}^P . Realizing that our simpler trading strategy does not discriminate between high or low expected trading profits, we compute P&L when trades occur only when an expected profit of at least $d = 15\%$ is expected. Now we observe that all models can be profitable even under 0.4% transaction costs.

Overall we may conclude the following. First, without transaction costs, the CS models produce higher P&L than their counterpart BS models. This is expected since the delta values generated by CS models are consistently higher than those of BS models (for example the median delta values of BS_{con}^P for *AggTs* is 0.632 whilst for CS_{con}^P is 0.697), making CS based trading more aggressive. Moreover, CS with $\{av\}$ and $\{avT\}$ volatility measures, outperforms significantly the equivalent BS models since it generates more than twice the number of trades; this may happen because unlike the BS models whose implied volatility changes more smoothly, CS models implied skewness and kurtosis can change more erratically. Secondly, and for the same reason, CS models under 0.2% or 0.4% transaction costs become inferior to their BS counterparts. Thirdly, from unreported calculations we have seen that as d increases we generally observe P&L to increase in a diminishing fashion indicating that there is an optimal d for maximizing trading profits. Finally, trading “in aggregate” positions leads to significant further savings on transaction costs.

[Tables 6-8, here]

In Table 8 we present results for the trading strategies based on ANNs (only for the hybrid models with time adjusted parameters). In general we observe similar results to those of the parametric models. Contrary though to the parametric OPMs, the ANNs offer significant improvement in the cases of less sophisticated parameter estimates. For example, BS_{av}^{Nh*} produces a P&L equal to 32,908 compared to a P&L equal to 14,088 in the case of BS_{av}^P . The best models provide profits in 77%-82% of transactions (detailed figures not reported for brevity) using both the full and the reduced dataset. Finally, in the presence of transaction costs the BS based hybrid model with contract specific volatility is not only the best performing ANN model, but also the overall best. A final observation is that the ability to generate profits even

under a considerable level of transaction costs (we do not report here, but the best strategies retained profitability even up to a level of 0.5% of transaction costs) provides some evidence of inefficiency in these options markets. Our study however is not intended to be a test of market efficiency.

6.1. Improving trading performance with the Chen and Johnson (1985) modified hedging approach

We now extend the trading strategies by utilizing with all models the improved hedging scheme suggested by Chen and Johnson (1985). This is a widely neglected (see Roon et al., 1998 for a rare exception in the use of parametric models) approach that deals with deriving hedge parameters under the assumption of mispriced options. According to this hedging scheme and when an option is mispriced, the delta hedge parameter, $\Delta_{m,t}^k$, should be derived in a different way. If a mispriced option has been identified, then the riskless hedge will not earn r , the riskless rate, but some other rate, r^* . Chen and Johnson obtain the expression for a European call option that is the same as BS presented in Equations 1, 1.a and 1.b, by replacing r with r^* . In order to derive the correct hedge ratio, Equation 1 must be solved numerically for r^* using the observed market price of c^{mrk} (like retrieving the implied interest rate). We implement this approach with the parametric BS and CS models, and the ANNs.

Finding the implied interest rate, r^* , for the case of BS or CS is a simple numerical task and we employ the repeated cubic interpolation technique according to Charalambous (1992). Finding the implied interest rate, r^* , for ANNs is a more involved task, since in the case of hybrid models we need to jointly optimize with respect to the interest rate input to the neural networks and to the interest rate in the parametric model that is used to create the hybrid target function; this introduces many jagged ridge regions in the optimization surface. Thus, in the case of hybrid ANNs we adopt a more computationally intensive methodology according to which we again use the cubic interpolation technique with ten different initial starting points.

After finding r^* for all models considered we rerun the trading strategies. Results for the parametric BS and CS models appear in Panel B of Tables 6 and 7. The most important observation is that before transaction costs are accounted for, in all BS models under consideration there is a slight (only) improvement in their profitability (P&L). Under aggregate 0.4% transaction costs and for $d = 15\%$, the improvement in BS_{60}^P is about 19%, in BS_{vix}^P is surprisingly about 164% and for the

more sophisticated BS_{con}^P model only 1.67%. We remind that BS_{vix}^P exhibited both, the poorest out-of-sample pricing performance and only a modest profitability (under 0.4% transaction costs) among the BS models. Under the adjusted deltas, this seems to be partly alleviated. Somewhat similar results we observe for the semi-parametric CS model. For both parametric models, the modified hedging approach under transaction costs gave the best results when using the average (not contract specific) parameters. In the case of ANNs (results unreported for brevity) and under no transaction costs, we also observe a slight tendency for increased performance, but the results are mixed. With transaction costs the technique was unable to improve the profitability of ANNs. The above observations refer to the full dataset, but they also hold for the reduced one (again not reported due to brevity).

A general observation for the use of the modified hedging approach in trading strategies is that it significantly improves trading performance when it is applied with OPM models under assumptions consistent with the assumptions under which this approach was developed. Thus, it performs well with the parametric models when either historical, or average implied parameters are used. The use of this approach did not reverse our previous findings about the best performing models when trading in the presence of transaction costs. Still, it demonstrated that simple models can be efficient alternatives to the more sophisticated and computationally intensive hybrid ANN methods.

6.2. Delta hedging

We have also considered hedging as a testing tool. Our results here coincide with previous literature – model ranking may differ if testing is based on hedging instead of pricing. Bakshi et al. (1997) compare alternative parametric models and state that the hedging-based ranking of the models is in sharp contrast with that obtained based on out-of-sample pricing. They also state that (delta-hedging) performance is virtually indistinguishable among models. Quite similar results are reported in papers where non-parametric methods were used, like Garcia and Gencay (2000), and Gencay and Qi (2001). Schittenkopf and Dorffner (2001) find the results (marginally) better for the parametric models, but practically indistinguishable. Hutchison et al. (1994) also report that the learning networks they use have a better hedging performance compared to BS but they find it difficult to infer which network type performs best. We attribute this difference of model ranking to the fact that models are usually optimized with respect to pricing. An exception is Carverhill and Cheuk (2003) who focus more on hedging performance by optimizing

with respect to the hedge parameters. Optimizing the “hedging performance” is beyond the scope of our paper. Furthermore, hedging performance is not a substitute for trading performance, since hedging tests fail to account for the difference between overpriced and underpriced options.

We have calculated the mean hedging error (MHE) and the mean absolute hedging error (MAHE) of a standard hedging strategy with daily rebalancing. For brevity we do not report the full results here, but we have found according to MHE that the best parametric model is the CS_{con}^P . Among the ANN models the best performing one is CS_{con}^{Nh*} , with an identical error for the parametric CS model (equal for both models to 0.26). In addition, the error equals 0.30 for both the BS_{con}^P and the BS_{con}^{Nh*} models. In general, from the MHE we cannot tell which OPM is the best since their difference in this measure is practically indistinguishable. Continuing with the MAHE we have the same picture, and we find it hard to observe a certain OPM that dominates in this measure since many models have “almost identical” MAHE values. It is true that BS_{con}^P and BS_{avT4}^P are the overall best models (with MAHE equal to 2.57 for both) and perform relatively better than the ANN models (their hybrid ANN counterparts both having an error equal to 2.63).

In general, we can conclude that the hedging error performance is not in line with the models’ pricing performance. That is, our best model in pricing accuracy, BS_{con}^{Nh*} , does not produce the smallest hedging errors. But again, it is truly hard to differentiate among models. The above discussion pertains to the full dataset, but we have observed that ranking models using hedging performance is not affected by the choice of dataset.

7. Conclusions

Our effort has focused in developing European option pricing and trading tools by combining the use of ANN methodology and information provided by parametric OPMs (the BS and the CS model). For our empirical tests we have used European call options on the S&P 500 Index from January 1998 to August 2001. In our analysis we have included historical parameters, a VIX volatility proxy derived by weighting implied volatilities (for the case of BS only), and implied parameters (an overall average, an average per maturity, the 4-point closest in moneyness, and a contract-specific parameter set). Neural networks are optimized using a modified Levenberg-Marquardt training algorithm. We include in the analysis simple ANNs (with input supplemented by historical or implied parameters specific either to BS or

the CS model), and hybrid ANNs that in addition use pricing information derived by any of the two parametric models. In order to check the robustness of the results, in addition to our *full* dataset we repeat the analysis using a *reduced* dataset (following Hutchison et al., 1994). The economic significance of the models is investigated through trading strategies with transaction costs. Instead of *naive* trading strategies we implement improved (dynamic and cost-effective) ones. Furthermore, we also refine these strategies with the Chen and Johnson (1985) modified hedging approach. Our results can be synopsized as follows:

Regarding the in-sample pricing, CS performs better than the BS model (with the exception of the case of the contract specific implied parameters that practically eliminate the pricing error).

Regarding out-of-sample pricing, CS outperforms BS with the use of average implied parameters, but BS is still a better model when the contract specific implied parameters are used; in general, implied parameters lead to better performance than the historical ones or the VIX volatility proxy; the simple neural networks cannot outperform the parametric models in the full range of data, but we verified allegations to the contrary found in the literature with the use of a reduced data set; hybrid neural networks that combine both neural network technology and the parametric models provide the best performance, especially when contract specific and adjusted parameters are used. The BS based hybrid ANN (with contract specific parameters) is the overall best performer, and the equivalent CS hybrid often a good alternative.

In trading and before transaction costs, models using contract specific implied parameters provide the best performance. But they also lead to the highest number of trades. In trading when transaction costs are accounted for in a *naive* manner, profits practically in all cases disappear. In trading and even with 0.4% transaction costs, when dynamic cost-efficient strategies are implemented, profits are still feasible hinting thus to potential market inefficiencies. The parametric BS with contract specific volatility is the best among the parametric models. The hybrid ANN based on BS with contract specific volatility is the overall best.

Implementing the widely neglected Chen and Johnson (1985) modified hedging approach, improves significantly the profitability of trading strategies that are based on the parametric models with average implied parameters (the models more consistent with the assumptions behind the modified hedging approach). This approach did not affect the choice of the overall best model in terms of trading with transaction costs. But it did demonstrate that reasonable alternatives for trading do exist without the need to resort to the extra sophistication of ANN technology.

References

- Andersen, T.G., Benzoni, L., Lund, J., 2002. An empirical investigation of continuous-time equity return models, *Journal of Finance* 57 (3) 1239-1276.
- Bakshi, G., Cao, C., Chen, Z., 1997. Empirical performance of alternative options pricing models, *Journal of Finance* 52 (5) 2003-2049.
- Bakshi, G., Cao, C., Chen, Z., 2000. Pricing and hedging long-term options, *Journal of Econometrics* 94 (1-2) 277-318.
- Bates, D.S., 1991. The Crash of '87: Was it expected? The evidence from options markets, *Journal of Finance* 46 (3) 1009-1044.
- Bates, D.S., 1996. Jumps and stochastic volatility: Exchange rate processes implicit in Deutsche mark options, *The Review of Financial Studies* 9 (1) 69-107.
- Bates, D.S., 2003. Empirical option pricing: A retrospection, *Journal of Econometrics* 116 (1-2) 387-404.
- Black, F., Scholes, M., 1972. The valuation of option contracts and a test of market efficiency, *The Journal of Finance* 27 (2) 399-417.
- Black, F., Scholes, M., 1973. The pricing of options and corporate liabilities, *Journal of Political Economy* 81 (3) 637-654.
- Black, F., Scholes, M., 1975. Fact and fantasy in the use of options, *The Financial Analysts Journal* 31, 36-41 and 61-72.
- Brown, C., Robinson, D., 2002. Skewness and kurtosis implied by option prices: A correction, *Journal of Financial Research* 25 (2) 279-282.
- Canica, L., Figlewski, S., 1993. The informational content of implied volatility, *The Review of Financial Studies* 6 (3) 659-681.
- Carverhill, A., Cheuk, T.H.F., 2003. Alternative neural network approach for option pricing and hedging, Working paper, School of Business, University of Hong Kong.
- Charalambous, C., 1992. A conjugate gradient algorithm for efficient training of artificial neural networks, *IEE Proceedings – G*, 139, 301-310.
- Chen, N., Johnson, H., 1985. Hedging options, *Journal of Financial Economics* 14 (2) 317-321.
- Cont, R., Fonseca, J., 2002. Dynamics of implied volatility surfaces, *Quantitative Finance* 2 (1) 45-60.
- Corrado, C.J., Su, T., 1996. Skewness and kurtosis in S&P 500 index returns implied by option prices, *Journal of Financial Research* 19 (2) 175-192.
- Cybenko, G., 1989. Approximation by superpositions of a sigmoidal function, *Mathematics of Control, Signal and Systems* 2, 303-314.
- Galai, D., 1977. Tests of market efficiency of the Chicago Board Options Exchange, *The Journal of Business* 50 (2) 167-197.

Garcia, R., Gencay, R., 2000. Pricing and hedging derivative securities with neural networks and a homogeneity hint, *Journal of Econometrics* 94 (1-2) 93-115.

Gencay, R., and Qi, M., 2001. Pricing and hedging derivative securities with neural networks: Bayesian regularization, early stopping and bagging, *IEEE Transactions on Neural Networks* 12 (4) 726-734.

Gultekin, N.B., Rogalski, R.J., Tinic S.M., 1982. Option pricing models estimates: some empirical results, *Financial Management* 11, 58-69.

Hagan, M., Demuth, H., Beale, M., 1996. *Neural Network Design*, PWS Publishing Company.

Hagan, M.T., Menhaj, M., 1994. Training feedforward networks with the Marquardt algorithm, *IEEE Transactions on Neural Networks* 5 (6) 989-993.

Hull, C.J., 1999. *Options Futures and Other Derivatives*, Prentice-Hall (4th edition).

Hutchison, J.M., Lo, A.W., Poggio, T., 1994. A nonparametric approach to pricing and hedging derivative securities via learning networks, *Journal of Finance* 49 (3) 851-889.

Lajbcygier P., Boek C., Palaniswami M., Flitman A., 1996. Comparing conventional and artificial neural network models for the pricing of options on futures, *Neurovest Journal* 4 (5) 16-24.

Lajbcygier, P., Flitman, A., Swan, A., Hyndman, R., 1997. The pricing and trading of options using a hybrid neural network model with historical volatility, *Neurovest Journal* 5 (1) 27-41.

Latane, H.A., Rendleman, R.J. Jr., 1976. Standard deviations of stock price ratios implied in option prices, *The Journal of Finance* 31 (2) 369-381.

MacBeth, J.D., Merville, L.J., 1980. Tests of the Black-Scholes and Cox call option valuation models, *Journal of Finance* 35 (2) 285-301.

Merton, R.C., 1973, Theory of rational option pricing, *Bell Journal of Economics and Management Science* 4 (Spring) 141-183.

Roon, De F., Veld C., Wei, J., 1998. A study on the efficiency of the market for Dutch long-term call options, *European Journal of Finance* 4 (2) 93-111.

Rubinstein, M., 1985. Nonparametric tests of alternative option pricing models using all reported trades and quotes on the 30 most active CBOE option classes from August 23, 1976 through August 31, 1978, *The Journal of Finance* 40 (2) 455-480.

Schittenkopf C., Dorffner, G., 2001. Risk-neutral density extraction from option prices: Improved pricing with mixture density networks, *IEEE Transactions on Neural Networks* 12 (4) 716-725.

Vahamaa, S., 2003. Skewness and kurtosis adjusted Black-Scholes model: A note on hedging performance, *Finance Letters* 1(5).

- Watson, P., Gupta, K.C., 1996. EM-ANN models for microstrip vias and interconnects in dataset circuits, *IEEE Transactions on Microwave Theory and Techniques* 44 (12) 2495-2503.
- Whaley, R.E., 1982. Valuation of American Call Options on dividend-paying stocks, *The Journal of Financial Economics* 10 (1) 29-58.
- Yao, J., Li, Y., Tan, C.L., 2000. Option price forecasting using neural networks, *The International Journal of Management Science* 28 (4) 455-466.

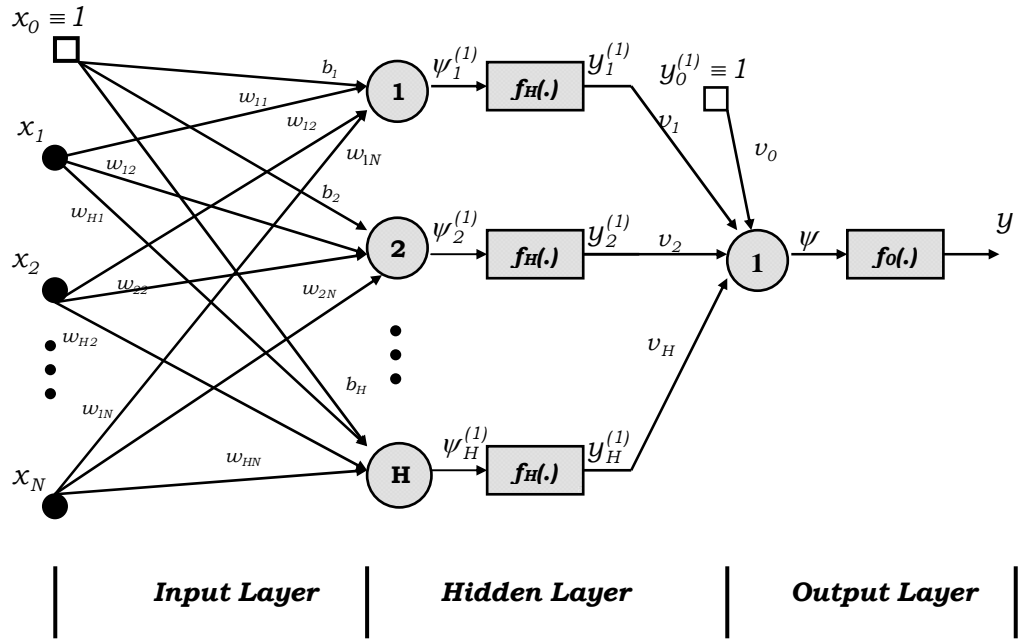


Fig. 1. A single hidden layer feedforward neural network

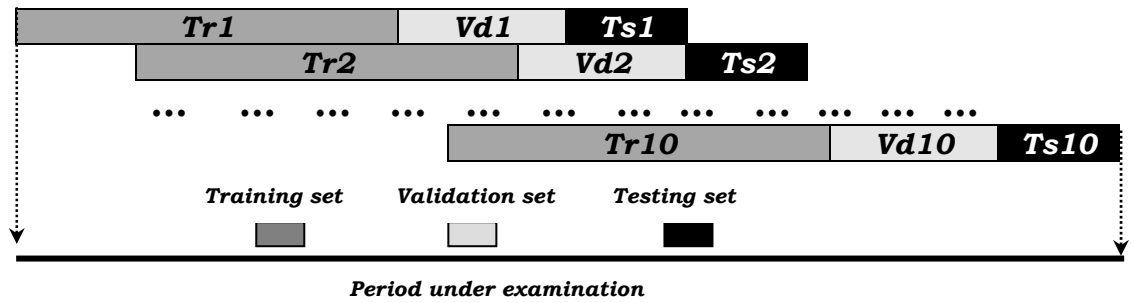


Fig. 2. The rolling-over training/validation/testing procedure

Table 1
Sample descriptive statistics

	VDOTM	DOTM	OTM	JOTM	ATM	JITM	ITM	DITM	VDITM
S/X	<0.85	0.85- 0.95	0.90- 0.95	0.95- 0.99	0.99- 1.01	1.01- 1.05	1.05- 1.10	1.10- 1.35	≥1.35
Short Term Options <60 Days									
Call	3.61	1.63	5.15	15.70	32.40	56.58	99.55	199.77	470.38
volatility	0.36	0.21	0.19	0.19	0.20	0.22	0.27	0.38	0.99
# obs	399	1,361	4,815	7,483	3,964	6,548	4,970	7,990	2,103
Medium Term Options 60-180 Days									
Call	4.38	8.29	23.58	46.06	64.51	90.35	131.10	227.41	493.18
volatility	0.22	0.18	0.20	0.21	0.21	0.23	0.25	0.30	0.54
# obs	1,412	1,727	2,578	3,147	1,780	2,901	3,038	8,100	3,999
Long Term Options ≥ 180 Days									
Call	9.65	42.09	74.03	106.24	126.03	150.99	185.87	267.12	495.82
Volatility	0.18	0.21	0.22	0.23	0.24	0.25	0.26	0.28	0.40
# obs	332	333	575	603	343	660	812	2,695	1,733

Sample characteristics for the period January 5, 1998 to August 24, 2001 concerning the average call option value, the average Black and Scholes contract specific implied volatility and the number of observations in each moneyness/maturity class.

Table 2

Pricing error measures in the aggregate testing period (*AggTs*)

	BS_{60}^P	BS_{vix}^P	BS_{av}^P	BS_{avT}^P	BS_{avT4}^P	BS_{con}^P	CS_{60}^P		CS_{av}^P	CS_{avT}^P	CS_{avT4}^P	CS_{con}^P	
RMSE	11.18	12.57	9.72	9.47	8.03	7.04	11.25		8.89	8.87	8.11	7.71	
MAE	6.83	8.60	5.32	5.00	3.10	2.70	6.89		3.86	3.72	3.27	3.10	
MeAE	4.48	6.38	3.74	3.37	1.52	1.43	4.61		2.26	1.94	1.69	1.68	
	BS_{60}^N	BS_{vix}^N	BS_{av}^N	BS_{avT}^N	BS_{avT4}^N	BS_{con}^N	BS_{60}^{N*}		BS_{vix}^{N*}	BS_{av}^{N*}	BS_{avT}^{N*}	BS_{avT4}^{N*}	BS_{con}^{N*}
RMSE	13.06	12.65	10.97	12.48	10.74	9.06	14.68		12.76	12.30	11.69	9.33	7.86
MAE	7.58	6.65	5.91	7.04	6.04	4.68	7.68		6.70	6.67	6.55	5.04	3.81
MeAE	5.13	3.83	3.65	4.11	3.69	2.88	4.71		3.65	3.99	3.94	2.94	2.44
	CS_{60}^N	CS_{av}^N		CS_{avT}^N	CS_{avT4}^N	CS_{con}^N	CS_{60}^{N*}		CS_{av}^{N*}	CS_{avT}^{N*}	CS_{avT4}^{N*}	CS_{con}^{N*}	
RMSE	15.22	11.28		11.59	9.87	11.83	14.35		11.42	11.96	9.47	9.76	
MAE	9.13	5.80		6.14	5.73	5.81	7.71		5.39	5.56	4.67	4.87	
MeAE	6.43	3.48		3.96	3.65	3.65	4.27		3.26	3.15	2.93	3.03	
	BS_{60}^{Nh}	BS_{vix}^{Nh}	BS_{av}^{Nh}	BS_{avT}^{Nh}	BS_{avT4}^{Nh}	BS_{con}^{Nh}	BS_{60}^{Nh*}	BS_{vix}^{Nh*}	BS_{av}^{Nh*}	BS_{avT}^{Nh*}	BS_{avT4}^{Nh*}	BS_{con}^{Nh*}	
RMSE	9.05	8.35	8.57	8.29	7.79	6.38	9.03	8.27	8.87	7.84	7.68	6.01	
MAE	5.40	4.55	4.35	4.09	3.30	2.68	5.46	4.53	4.35	3.91	3.17	2.61	
MeAE	3.73	2.98	2.83	2.51	1.80	1.60	3.98	3.00	2.69	2.53	1.67	1.58	
	CS_{60}^{Nh}	CS_{av}^{Nh}		CS_{avT}^{Nh}	CS_{avT4}^{Nh}	CS_{con}^{Nh}	CS_{60}^{Nh*}		CS_{av}^{Nh*}	CS_{avT}^{Nh*}	CS_{avT4}^{Nh*}	CS_{con}^{Nh*}	
RMSE	10.33	8.68		8.63	7.97	7.60	9.68		8.83	8.66	7.60	7.39	
MAE	6.38	4.12		3.84	3.42	3.14	6.20		3.95	3.94	3.39	3.11	
MeAE	4.46	2.42		2.17	1.93	1.77	4.56		2.33	2.35	1.96	1.82	

RMSE is the Root Mean Square Error, MAE the Mean Absolute Deviation and MeAE the Median of the Absolute Error. The superscripts refer to the kind of the model: *P* refers to parametric models, *N* to the simple neural networks and *Nh* to the hybrid neural networks. The asterisk (*) refers to neural network models that use transformed variables. The subscripts refer to the kind of historical/implied parameters used.

Table 3
Matched-pair student t -tests for square and absolute differences

	BS_{60}^P	BS_{vix}^P	BS_{con}^P	CS_{60}^P	CS_{con}^P	BS_{60}^{N*}	CS_{60}^{N*}	BS_{60}^{Nh*}	BS_{vix}^{Nh*}	BS_{con}^{Nh*}	CS_{60}^{Nh*}	CS_{con}^{Nh*}
BS_{60}^P		-27.74	75.12	-0.94	65.72	-11.07	-11.72	23.90	40.83	81.22	10.80	66.92
BS_{vix}^P	7.17		104.84	26.74	94.84	11.84	11.71	53.72	70.71	112.32	40.52	96.53
BS_{con}^P	-16.13	-25.08		-75.91	-8.43	-70.51	-72.82	-56.94	-38.56	2.12	-70.87	-8.76
CS_{60}^P	0.34	-6.72	16.31		66.53	-10.28	-10.91	24.85	41.74	82.02	11.78	67.74
CS_{con}^P	-13.38	-21.60	2.14	-13.58		-63.58	-65.63	-46.76	-28.85	11.11	-60.38	-0.10
BS_{60}^{N*}	7.24	4.64	13.37	7.09	12.48		-0.34	30.64	43.94	74.23	20.23	64.28
CS_{60}^{N*}	7.77	4.67	15.19	7.59	14.09	-0.62		31.84	45.50	76.80	21.15	66.39
BS_{60}^{Nh*}	-9.55	-18.30	7.57	-9.81	4.95	-10.83	-12.15		18.64	63.30	-14.28	47.82
BS_{vix}^{Nh*}	-12.54	-21.61	4.48	-12.75	2.02	-11.91	-13.46	-3.25		43.70	-32.87	29.50
BS_{con}^{Nh*}	-21.16	-32.03	-3.45	-21.26	-5.62	-14.65	-16.83	-12.27	-8.78		-78.03	-11.58
CS_{60}^{Nh*}	-6.86	-15.36	10.42	-7.15	7.65	-9.84	-10.96	2.97	6.24	15.52		61.73
CS_{con}^{Nh*}	-14.98	-23.78	1.15	-15.16	-1.04	-12.95	-14.69	-6.34	-3.26	4.73	-9.18	

Reported matched-pair t -tests concerning the absolute differences are in the upper diagonal, whilst the matched-pair t -tests concerning the square differences in the lower diagonal. Both tests compare the MAE and MSE between models in the vertical heading versus models in the horizontal heading. In general, a positive t -value larger than 1.645 (2.325) means that the model in the vertical heading has a larger MAE or MSE than the model in the horizontal heading at 5% (1%) significance level.

Table 4
Pricing error measures in the aggregate testing period (*AggTs*) for the reduced dataset

	BS_{60}^P	BS_{vix}^P	BS_{av}^P	BS_{avT}^P	BS_{avT4}^P	BS_{con}^P	CS_{60}^P		CS_{av}^P	CS_{avT}^P	CS_{avT4}^P	CS_{con}^P	
RMSE	9.83	11.82	8.41	8.25	7.08	7.06	9.74		7.56	7.55	7.55	7.52	
MAE	6.35	8.43	4.82	4.54	2.65	2.65	6.32		3.38	3.12	2.99	3.04	
MeAE	4.50	6.57	3.63	3.27	1.48	1.46	4.59		2.17	1.83	1.69	1.71	
	BS_{60}^N	BS_{vix}^N	BS_{av}^N	BS_{avT}^N	BS_{avT4}^N	BS_{con}^N	BS_{60}^{N*}		BS_{vix}^{N*}	BS_{av}^{N*}	BS_{avT}^{N*}	BS_{avT4}^{N*}	BS_{con}^{N*}
RMSE	8.05	6.56	7.34	6.94	6.64	6.69	7.14		6.60	6.82	6.91	6.25	6.12
MAE	5.07	3.34	4.02	3.72	3.42	3.37	4.11		3.43	3.46	3.59	3.01	3.00
MeAE	3.80	2.32	2.99	2.56	2.33	2.24	3.09		2.41	2.44	2.56	1.99	2.02
	CS_{60}^N	CS_{av}^N		CS_{avT}^N	CS_{avT4}^N	CS_{con}^N	CS_{60}^{N*}		CS_{av}^{N*}	CS_{avT}^{N*}	CS_{avT4}^{N*}	CS_{con}^{N*}	
RMSE	9.05	7.18		6.93	6.94	6.88	8.35		6.97	6.59	6.50	6.77	
MAE	5.74	3.95		3.61	3.73	3.62	4.94		3.68	3.26	3.23	3.45	
MeAE	4.25	2.74		2.41	2.60	2.55	3.43		2.62	2.22	2.25	2.36	
	BS_{60}^{Nh}	BS_{vix}^{Nh}	BS_{av}^{Nh}	BS_{avT}^{Nh}	BS_{avT4}^{Nh}	BS_{con}^{Nh}	BS_{60}^{Nh*}	BS_{vix}^{Nh*}	BS_{av}^{Nh*}	BS_{avT}^{Nh*}	BS_{avT4}^{Nh*}	BS_{con}^{Nh*}	
RMSE	8.45	6.70	7.29	7.01	6.58	6.78	7.35	6.40	7.05	6.83	5.94	5.64	
MAE	5.11	3.58	3.62	3.38	2.62	2.69	4.27	3.21	3.32	3.30	2.45	2.44	
MeAE	3.44	2.59	2.55	2.35	1.55	1.65	3.13	2.26	2.30	2.33	1.51	1.54	
	CS_{60}^{Nh}	CS_{av}^{Nh}		CS_{avT}^{Nh}	CS_{avT4}^{Nh}	CS_{con}^{Nh}	CS_{60}^{Nh*}		CS_{av}^{Nh*}	CS_{avT}^{Nh*}	CS_{avT4}^{Nh*}	CS_{con}^{Nh*}	
RMSE	7.80	7.29		6.83	7.31	7.35	7.69		6.90	6.80	6.51	6.46	
MAE	4.65	3.20		3.08	3.03	3.03	4.58		3.13	2.92	2.83	2.87	
MeAE	3.41	2.13		2.02	1.82	1.80	3.23		2.03	1.80	1.79	1.81	

RMSE is the Root Mean Square Error, MAE the Mean Absolute Error and MeAE the Median of the Absolute Error. The superscripts refer to the kind of the model: *P* refers to parametric models, *N* to the simple neural networks and *Nh* to the hybrid neural networks. The asterisk (*) refers to neural network models that use the transformed variables. The subscripts refer to the kind of historical/implied parameters used.

Table 5

Root Mean Square Errors for selected models (clustered by moneyness and maturity)

	Short	Medium	Long	Short	Medium	Long
Results for the full dataset						
	BS_{con}^P			CS_{con}^P		
VDOTM	3.60	4.91	0.56	8.34	10.61	0.66
DOTM	2.27	4.50	2.82	3.02	5.24	4.47
OTM	5.78	8.37	3.97	6.29	9.68	5.08
JOTM	7.81	6.68	6.15	8.13	7.64	7.65
ATM	6.67	9.46	5.86	7.30	10.14	7.29
JITM	6.71	9.41	4.34	7.29	9.21	5.97
ITM	7.70	7.13	4.43	8.24	7.59	5.18
DITM	7.07	7.93	7.27	7.20	8.50	7.50
VDITM	8.26	9.46	8.74	8.29	10.05	9.05
	BS_{con}^{Nh*}			CS_{con}^{Nh*}		
VDOTM	3.60	4.97	1.15	6.13	10.22	6.04
DOTM	2.46	4.83	2.32	2.96	5.28	5.03
OTM	5.50	7.75	3.98	6.19	9.41	5.36
JOTM	5.89	5.36	5.78	7.83	7.30	7.66
ATM	4.73	8.18	5.38	6.94	9.86	7.13
JITM	5.59	7.39	4.10	6.89	8.68	6.64
ITM	6.24	6.05	3.95	7.58	7.16	5.69
DITM	5.80	7.15	6.74	6.64	8.04	7.17
VDITM	8.03	9.29	8.46	8.96	10.33	9.26
Results for the reduced dataset						
	BS_{con}^{Nh*}			CS_{con}^{Nh*}		
VDOTM	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
DOTM	2.36	4.07	n.a.	2.54	5.22	n.a.
OTM	5.08	7.25	n.a.	5.69	8.74	n.a.
JOTM	5.82	5.59	n.a.	6.76	7.09	n.a.
ATM	4.65	8.37	n.a.	5.68	9.53	n.a.
JITM	5.50	7.68	n.a.	6.20	8.16	n.a.
ITM	5.98	5.84	n.a.	6.73	6.75	n.a.
DITM	5.45	6.59	n.a.	5.95	7.67	n.a.
VDITM	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.

Table 6
Trading strategies for the Black and Scholes models

	BS_{60}^P	BS_{vix}^P	BS_{av}^P	BS_{avT}^P	BS_{avT4}^P	BS_{con}^P
Panel A: Black and Scholes trading strategy with standard delta values						
P&L	7,447	13,518	14,088	13,069	32,040	35,026
# Trades	3,361	3,878	4,858	5,477	13,539	15,644
P&L 0.2% ($d=0\%$)	-6,829	-6,847	-5,348	-7,512	-17,911	-23,307
Agg P&L 0.2% ($d=0\%$)	-1,861	-266	737	-1,394	-5,638	-8,437
P&L 0.2% ($d=15\%$)	3,320	4,134	7,527	6,841	7,907	7,369
Agg P&L 0.2% ($d=15\%$)	5,003	5,019	8,344	7,657	8,384	7,873
P&L 0.4% ($d=0\%$)	-21,105	-27,211	-24,785	-28,093	-67,863	-81,640
Agg P&L 0.4% ($d=0\%$)	-11,170	-14,049	-12,614	-15,858	-43,316	-51,899
P&L 0.4% ($d=15\%$)	-1,468	-508	3,241	2,269	4,691	4,212
Agg P&L 0.4% ($d=15\%$)	1,897	1,262	4,875	3,901	5,645	5,221
Panel B: Black and Scholes trading strategy with modified delta values						
P&L	7,916	14,367	14,232	13,441	32,281	35,229
# Trades	-6,169	-5,599	-4,958	-6,946	-17,788	-23,080
P&L 0.2% ($d=0\%$)	-1,392	1,342	1,225	-778	-5,534	-8,259
Agg P&L 0.2% ($d=0\%$)	4,044	5,534	8,182	7,546	8,306	7,713
P&L 0.2% ($d=15\%$)	5,515	6,558	9,115	8,524	8,815	8,198
Agg P&L 0.2% ($d=15\%$)	-20,254	-25,564	-24,148	-27,334	-67,858	-81,390
P&L 0.4% ($d=0\%$)	-10,700	-11,682	-11,782	-14,998	-43,348	-51,748
Agg P&L 0.4% ($d=0\%$)	-685	1,284	4,143	3,180	4,883	4,339
P&L 0.4% ($d=15\%$)	2,257	3,333	6,007	5,137	5,900	5,308

P&L is the total profit and loss without transaction costs; # Trades is the number of trades. P&L ($d=0$ and 15%) represents the P&L at 0.2% or 0.4% transaction costs when we ignore trades whose absolute percentage of mispricing error between model estimates and market values is at least 0% and 15% respectively. *Agg.* refers to aggregating the position on the underlying asset to reduce transaction costs. Panel A tabulates results with standard delta values whilst Panel B tabulates results with Chen and Johnson modified delta values.

Table 7
Trading strategies for the Corrado and Su models

	CS_{60}^P	CS_{av}^P	CS_{avT}^P	CS_{avT4}^P	CS_{con}^P
Panel A: Corrado and Su trading strategy with standard delta values					
P&L	7,603	28,816	32,803	37,072	36,777
# Trades	3,430	11,178	13,306	14,911	15,219
P&L 0.2% ($d=0\%$)	-7,658	-15,867	-19,045	-22,750	-24,414
Agg P&L 0.2% ($d=0\%$)	-2,532	-4,495	-5,641	-6,685	-6,909
P&L 0.2% ($d=15\%$)	2,868	7,960	6,791	6,606	6,422
Agg P&L 0.2% ($d=15\%$)	4,533	8,739	7,483	7,418	7,311
P&L 0.4% ($d=0\%$)	-22,919	-60,550	-70,894	-82,572	-85,604
Agg P&L 0.4% ($d=0\%$)	-12,667	-37,805	-44,085	-50,441	-50,595
P&L 0.4% ($d=15\%$)	-1,949	2,797	1,935	1,371	1,124
Agg P&L 0.4% ($d=15\%$)	1,383	4,355	3,319	2,993	2,901
Panel B: Corrado and Su trading strategy with modified delta values					
P&L	7,837	29,208	33,219	37,044	37,097
# Trades	3,430	11,178	13,306	14,911	15,219
P&L 0.2% ($d=0\%$)	-7,209	-15,317	-18,610	-22,828	-24,203
Agg P&L 0.2% ($d=0\%$)	-2,332	-3,843	-5,186	-6,708	-6,615
P&L 0.2% ($d=15\%$)	3,512	8,685	7,322	6,740	6,778
Agg P&L 0.2% ($d=15\%$)	4,943	9,539	8,024	7,594	7,720
P&L 0.4% ($d=0\%$)	-22,255	-59,841	-70,439	-82,700	-85,503
Agg P&L 0.4% ($d=0\%$)	-12,501	-36,893	-43,590	-50,460	-50,328
P&L 0.4% ($d=15\%$)	-1,218	3,521	2,303	1,172	1,074
Agg P&L 0.4% ($d=15\%$)	1,646	5,229	3,707	2,881	2,958

P&L is the total profit and loss without transaction costs; # Trades is the number of trades. P&L ($d=0$ and 15%) represents the P&L at 0.2% or 0.4% transaction costs when we ignore trades whose absolute percentage of mispricing error between model estimates and market values is at least 0% and 15% respectively. *Agg.* refers to aggregating the position on the underlying asset to reduce transaction costs. Panel A tabulates results with standard delta values whilst Panel B tabulates results with Chen and Johnson modified delta values.

Table 8
Trading strategies for the hybrid ANN models

	BS_{60}^{Nh*}	BS_{vix}^{Nh*}	BS_{av}^{Nh*}	BS_{avT}^{Nh*}	BS_{avT4}^{Nh*}	BS_{con}^{Nh*}
Panel A: Black and Scholes based hybrid ANNs						
P&L	27,024	29,529	32,908	33,514	35,774	37,281
# Trades	5,675	8,246	8,907	9,457	11,995	12,650
P&L 0.2% ($d=0\%$)	1,694	-4,193	-2,435	-4,134	-11,484	-12,939
Agg P&L 0.2% ($d=0\%$)	10,552	6,053	7,871	7,086	837	1,066
P&L 0.2% ($d=15\%$)	6,593	5,147	8,162	8,579	7,910	8,427
Agg P&L 0.2% ($d=15\%$)	8,247	6,977	9,890	9,957	8,689	9,237
P&L 0.4% ($d=0\%$)	-23,637	-37,914	-37,778	-41,782	-58,741	-63,158
Agg P&L 0.4% ($d=0\%$)	-5,920	-17,424	-17,166	-19,343	-34,100	-35,148
P&L 0.4% ($d=15\%$)	1,804	-277	2,232	3,156	4,364	4,812
Agg P&L 0.4% ($d=15\%$)	5,112	3,382	5,687	5,911	5,922	6,432
	CS_{60}^{Nh*}	CS_{av}^{Nh*}	CS_{avT}^{Nh*}	CS_{avT4}^{Nh*}	CS_{con}^{Nh*}	
Panel B: Corrado and Su based hybrid ANNs						
P&L	26,691		32,915	31,943	34,907	37,975
# Trades	5,140		10,043	10,377	12,537	12,947
P&L 0.2% ($d=0\%$)	3,590		-8,721	-12,019	-17,527	-16,084
Agg P&L 0.2% ($d=0\%$)	11,032		3,734	898	-1,586	735
P&L 0.2% ($d=15\%$)	7,337		6,653	5,601	6,052	7,826
Agg P&L 0.2% ($d=15\%$)	8,861		8,231	7,114	7,439	8,960
P&L 0.4% ($d=0\%$)	-19,511		-50,356	-55,980	-69,962	-70,143
Agg P&L 0.4% ($d=0\%$)	-4,626		-25,446	-30,146	-38,078	-36,505
P&L 0.4% ($d=15\%$)	2,433		724	457	612	2,605
Agg P&L 0.4% ($d=15\%$)	5,481		3,879	3,484	3,387	4,873

P&L is the total profit and loss without transaction costs; # Trades is the number of trades. P&L ($d=0$ and 15%) represents the P&L at 0.2% or 0.4% transaction costs when we ignore trades whose absolute percentage of mispricing error between model estimates and market values is at least 0% and 15% respectively. Agg. refers to aggregating the position on the underlying asset to reduce transaction costs. Panel A tabulates results for the hybrid BS based ANN model whilst Panel B tabulates results for the hybrid CS based ANN models.